

Sentiment Classification of Restaurant Reviews Using Machine Learning Techniques

Group 8 ML1

Import required libraries

```
In [2]: # Data handling
import pandas as pd
import numpy as np

# Data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Text preprocessing
import re
import string
from sklearn.feature_extraction.text import TfidfVectorizer

# Dimensionality reduction
from sklearn.decomposition import PCA

# Data preprocessing / model prep
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Machine Learning models
from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier

# Evaluation / metrics
from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, classification_report
import math
```

```
In [3]: # Load data
df = pd.read_csv("restaurant_reviews_sample.csv")
df.head(5)
```

Out[3]:

	review_id	user_id	business_id	stars
0	KU_O5udG6zpxOg-VcAEodg	mh_-eMZ6K5RLWhZyISBhwA	XQfwVwDr-v0ZS3_CbbE5Xw	3.0
1	saUsX_uimxRICVr67Z4Jig	8g_iMtfSiwikVnbP2etR0A	YjUWPPl6HXG530lwP-fb2A	3.0
2	AqPFMleE6RsU23_auESxiA	_7bHUi9Uuf5__HHc_Q8guQ	kxX2SOes4o-D3ZQBkiMRfA	5.0
3	Sx8TMOWLNUJBWer-0pcmoA	bcjbaE6dDog4jkNY91ncLQ	e4Vwtrqf-wpJfwesgvdgxQ	4.0
4	JrlxIS1TzJ-iCu79ul40cQ	eUta8W_HdHMXPzLBBZhL1A	04UD14gamNjLY0IDYVhHJg	1.0



Data preparation

In [4]:

```
# Keep required columns
df = df[['review_id', 'text', 'stars']].copy()

# Drop missing
df = df.dropna(subset=['text', 'stars'])

# Remove duplicate rows & duplicate IDs
df = df.drop_duplicates().drop_duplicates(subset='review_id')

# Rename the column 'text' to 'reviews'
df = df.rename(columns={'text': 'reviews'})
```

```

# Create sentiment label
def label_sentiment(star):
    if star >= 4:
        return 1
    elif star <= 2:
        return 0
    else:
        return None

df['sentiment'] = df['stars'].apply(label_sentiment)

# Drop neutral reviews
df = df.dropna(subset=['sentiment'])

# Keep only needed columns for modeling
reviews_df = df[['reviews', 'sentiment']].copy()

reviews_df.head()

```

Out[4]:

	reviews	sentiment
2	Wow! Yummy, different, delicious. Our favo...	1.0
3	Cute interior and owner (?) gave us tour of up...	1.0
4	I am a long term frequent customer of this est...	0.0
5	Amazingly amazing wings and homemade bleu chee...	1.0
7	Locals recommended Milktooth, and it's an amaz...	1.0

EDA

In [5]:

```

df = reviews_df.copy()

# Dataset Overview
df.head()
df.info()
df.describe(include='object')
df.isnull().sum()

<class 'pandas.core.frame.DataFrame'>
Index: 43720 entries, 2 to 49999
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype  
---  --          -----          ---  
 0   reviews     43720 non-null   object 
 1   sentiment   43720 non-null   float64 
dtypes: float64(1), object(1)
memory usage: 1.0+ MB

```

```

Out[5]: reviews      0
         sentiment    0
         dtype: int64

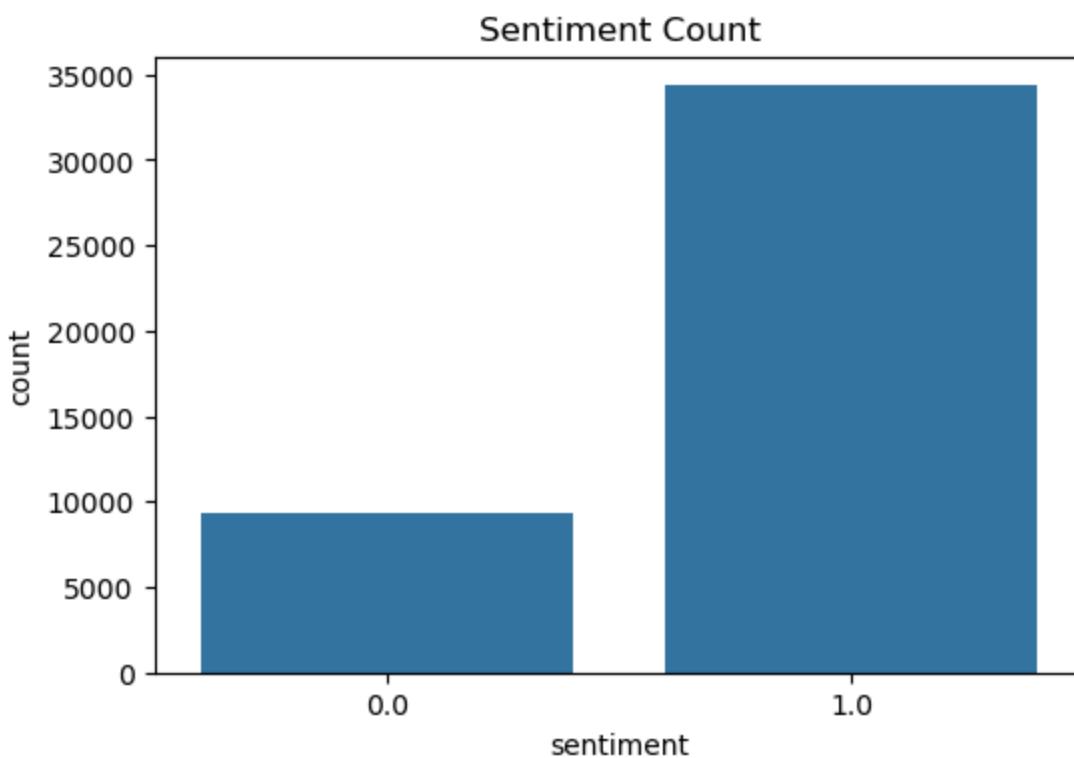
```

```
In [6]: #Counts positives and negatives  
df['sentiment'].value_counts()
```

```
Out[6]: sentiment  
1.0    34335  
0.0     9385  
Name: count, dtype: int64
```

Sentiment distribution

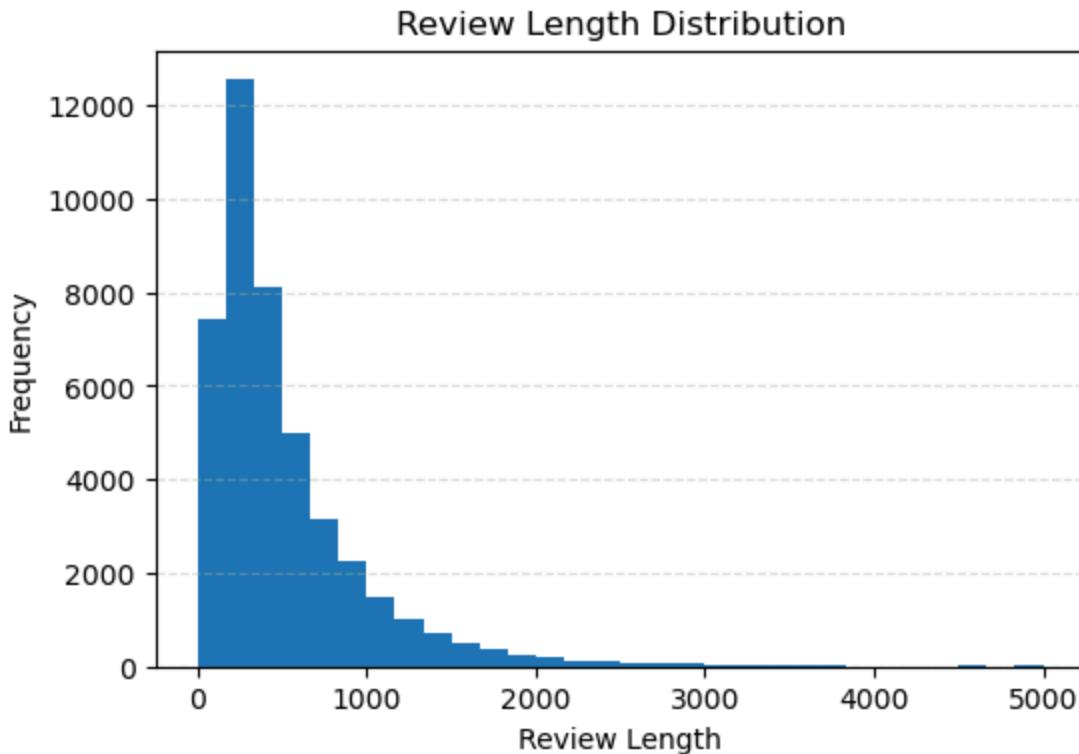
```
In [7]: #Sentiment distribution countplot  
plt.figure(figsize=(6,4))  
sns.countplot(x='sentiment', data=df)  
plt.title("Sentiment Count")  
plt.show()
```



Review length distribution

```
In [8]: #Creating review_length column  
df['review_length'] = df['reviews'].astype(str).apply(len)
```

```
In [9]: #Histogram of Review Length (distribution)  
plt.figure(figsize=(6,4))  
plt.hist(df['review_length'], bins=30)  
plt.xlabel("Review Length")  
plt.ylabel("Frequency")  
plt.title("Review Length Distribution")  
plt.grid(axis='y', linestyle='--', alpha=0.4)  
plt.show()
```



Basic statistics (mean, median, mode, etc.)

```
In [10]: #Summary Statistics (Mean, Median, Mode, Min, Max, SD)
print("Mean:", df['review_length'].mean())
print("Median:", df['review_length'].median())
print("Mode:", df['review_length'].mode()[0])
print("Std Dev:", df['review_length'].std())
print("Min:", df['review_length'].min())
print("Max:", df['review_length'].max())
```

Mean: 522.4263037511437
 Median: 369.0
 Mode: 233
 Std Dev: 486.472346504319
 Min: 3
 Max: 5000

Text Pre-Processing

```
In [11]: # Text Cleaning
import re
import string
from sklearn.feature_extraction.text import ENGLISH_STOP_WORDS

def clean_text(text):
    text = text.lower() # Lowercase
    text = re.sub(f"[{string.punctuation}]", "", text) # remove punctuation
    words = text.split()
    words = [w for w in words if w not in ENGLISH_STOP_WORDS] # remove stopwords
    return " ".join(words)
```

```
reviews_df['clean_reviews'] = reviews_df['reviews'].apply(clean_text)
```

Train/Test Split + TF-IDF Vectorization

```
In [12]: from sklearn.model_selection import train_test_split

# Define features and target
X = reviews_df['clean_reviews'] # cleaned text column
y = reviews_df['sentiment'] # target column (0/1)

# Split into train and test
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=1, stratify=y
)

# Check
print("Train size:", X_train.shape[0])
print("Test size:", X_test.shape[0])
```

Train size: 34976

Test size: 8744

```
In [14]: from sklearn.feature_extraction.text import TfidfVectorizer

# TF-IDF Vectorization
tfidf = TfidfVectorizer(
    max_features=5000, # limit features
    stop_words='english', # remove stopwords (optional, already cleaned)
    ngram_range=(1,2) # unigrams + bigrams
)

# Fit TF-IDF only on training data
X_train_tfidf = tfidf.fit_transform(X_train)

# Transform test data using the same TF-IDF
X_test_tfidf = tfidf.transform(X_test)

print("TF-IDF Train:", X_train_tfidf.shape)
print("TF-IDF Test :", X_test_tfidf.shape)
```

TF-IDF Train: (34976, 5000)

TF-IDF Test : (8744, 5000)

```
In [15]: from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (
    confusion_matrix,
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    classification_report
)
import pandas as pd
```

PCA + Scaling

```
In [17]: from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler

def apply_pca_and_scale(X_train, X_test, n=100):
    pca = PCA(n_components=n, random_state=1)
    X_train_pca = pca.fit_transform(X_train.toarray())
    X_test_pca = pca.transform(X_test.toarray())

    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train_pca)
    X_test_scaled = scaler.transform(X_test_pca)

    return X_train_scaled, X_test_scaled

X_train_lr, X_test_lr = apply_pca_and_scale(X_train_tfidf, X_test_tfidf)
X_train_knn, X_test_knn = apply_pca_and_scale(X_train_tfidf, X_test_tfidf)
X_train_mlr, X_test_mlr = apply_pca_and_scale(X_train_tfidf, X_test_tfidf)
```

```
In [21]: def print_classification_metrics(y_true, y_pred, model_name="Model"):
    cm = confusion_matrix(y_true, y_pred)
    acc = accuracy_score(y_true, y_pred)
    prec = precision_score(y_true, y_pred)
    rec = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)

    print(f"\n===== {model_name} =====")
    print("Confusion Matrix (rows = true, cols = predicted):")
    print(cm)
    print(f"\nAccuracy : {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall   : {rec:.4f}")
    print(f"F1-score : {f1:.4f}")
    print("\nDetailed classification report:")
    print(classification_report(y_true, y_pred, digits=4))
```

Models, Logistic regression

```
In [24]: from sklearn.linear_model import LogisticRegression

# Initialize model with class balancing
log_reg = LogisticRegression(
    solver='lbfgs',
    max_iter=1000,
    random_state=1,
    n_jobs=-1,
    class_weight='balanced'
)

log_reg.fit(X_train_lr, y_train)

# Predictions on test set
```

```
y_pred_lr = log_reg.predict(X_test_lr)

# Metrics for Logistic Regression
print_classification_metrics(y_test, y_pred_lr,
                             model_name="Logistic Regression")
```

===== Logistic Regression =====

Confusion Matrix (rows = true, cols = predicted):

```
[[1731 146]
 [ 655 6212]]
```

Accuracy : 0.9084

Precision: 0.9770

Recall : 0.9046

F1-score : 0.9394

Detailed classification report:

	precision	recall	f1-score	support
0.0	0.7255	0.9222	0.8121	1877
1.0	0.9770	0.9046	0.9394	6867
accuracy			0.9084	8744
macro avg	0.8513	0.9134	0.8758	8744
weighted avg	0.9230	0.9084	0.9121	8744

In [27]:

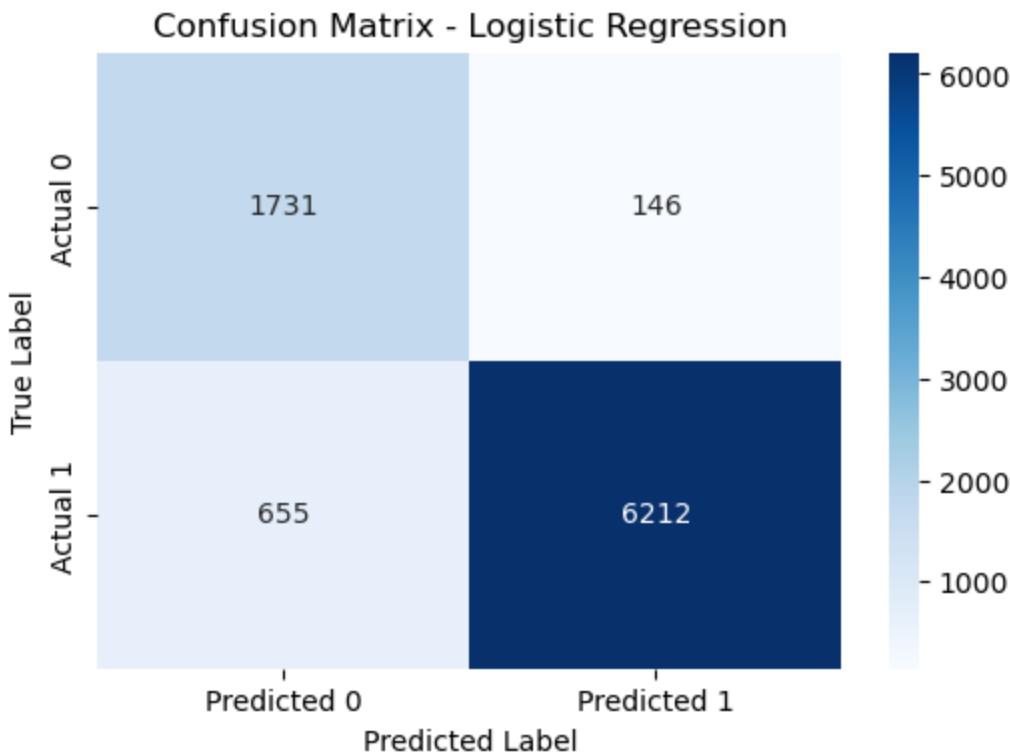
```
y_pred_lr = log_reg.predict(X_test_lr)
acc_lr = accuracy_score(y_test, y_pred_lr)
f1_lr = f1_score(y_test, y_pred_lr)
```

In [28]:

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

# Confusion matrix
cm = confusion_matrix(y_test, y_pred_lr)

plt.figure(figsize=(6,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.title("Confusion Matrix - Logistic Regression")
plt.ylabel("True Label")
plt.xlabel("Predicted Label")
plt.show()
```



KNN Classification model + metrics

```
In [32]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

# Try different k values and record accuracy
results = []
for k in range(1, 16):
    knn_tmp = KNeighborsClassifier(n_neighbors=k, weights='distance', n_jobs=-1)
    knn_tmp.fit(X_train_knn, y_train)
    y_pred_tmp = knn_tmp.predict(X_test_knn)
    acc_tmp = accuracy_score(y_test, y_pred_tmp)
    results.append({"k": k, "accuracy": acc_tmp})

results_df = pd.DataFrame(results)
print("KNN accuracy for different k values:")
print(results_df)

# Choose the k with highest accuracy on the test set
best_row = results_df.loc[results_df["accuracy"].idxmax()]
best_k = int(best_row["k"])
print(f"\nBest k based on test accuracy: k = {best_k}, "
      f"accuracy = {best_row['accuracy']:.4f}")

# Train final KNN model with best k
knn_best = KNeighborsClassifier(n_neighbors=best_k, weights='distance', n_jobs=-1)
knn_best.fit(X_train_knn, y_train)

# Predictions and metrics
y_pred_knn = knn_best.predict(X_test_knn)
```

```
print_classification_metrics(y_test, y_pred_knn,
                             model_name=f'KNN (k = {best_k}, PCA + scaled)')
```

KNN accuracy for different k values:

k	accuracy
0	0.743367
1	0.743481
2	0.793573
3	0.792086
4	0.810499
5	0.816560
6	0.824909
7	0.825595
8	0.827996
9	0.832113
10	0.836574
11	0.838289
12	0.842521
13	0.844122
14	0.846523

Best k based on test accuracy: k = 15, accuracy = 0.8465

===== KNN (k = 15, PCA + scaled) =====

Confusion Matrix (rows = true, cols = predicted):

```
[[1223 654]
 [ 688 6179]]
```

Accuracy : 0.8465
Precision: 0.9043
Recall : 0.8998
F1-score : 0.9020

Detailed classification report:

	precision	recall	f1-score	support
0.0	0.6400	0.6516	0.6457	1877
1.0	0.9043	0.8998	0.9020	6867
accuracy			0.8465	8744
macro avg	0.7721	0.7757	0.7739	8744
weighted avg	0.8476	0.8465	0.8470	8744

In [33]:

```
y_pred_knn = knn_best.predict(X_test_knn)
acc_knn = accuracy_score(y_test, y_pred_knn)
f1_knn = f1_score(y_test, y_pred_knn)
```

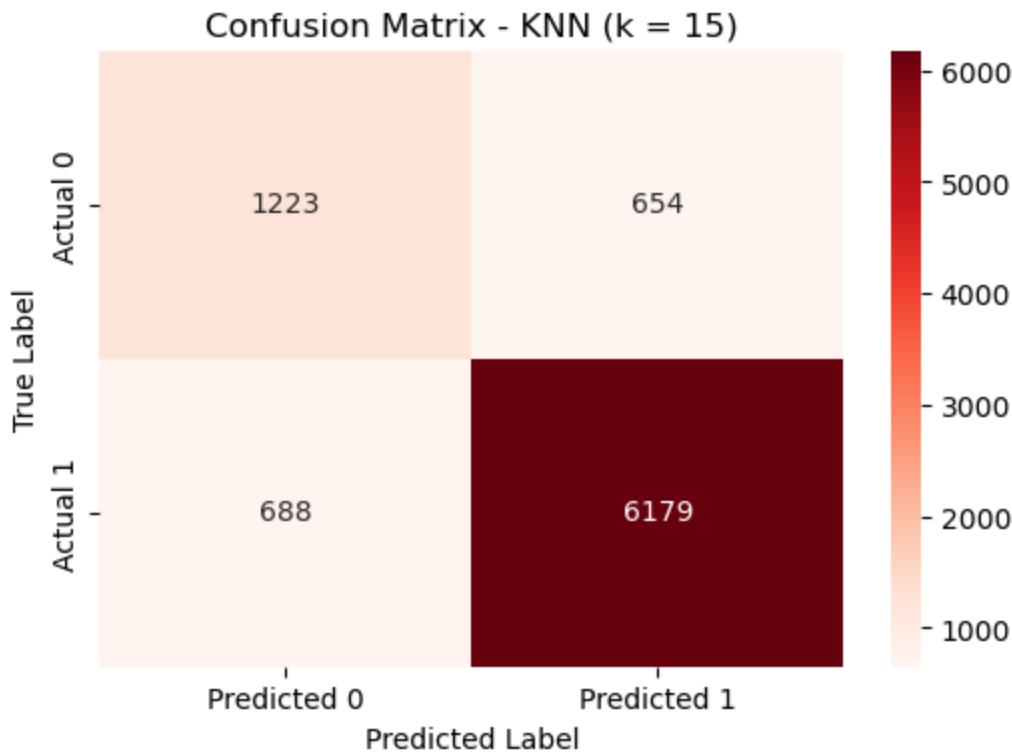
In [34]:

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

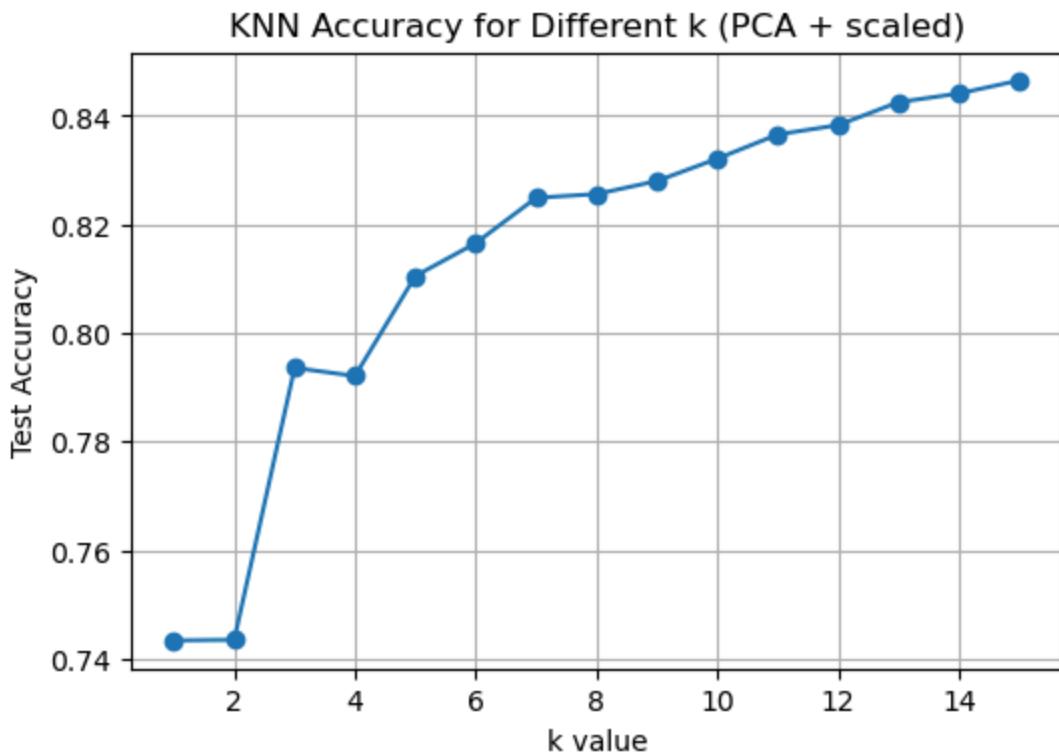
# Confusion matrix
cm_knn = confusion_matrix(y_test, y_pred_knn)

plt.figure(figsize=(6,4))
```

```
sns.heatmap(cm_knn, annot=True, fmt='d', cmap='Reds',
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.title(f"Confusion Matrix - KNN (k = {best_k})")
plt.ylabel("True Label")
plt.xlabel("Predicted Label")
plt.show()
```



```
In [35]: # Plot accuracy vs k
import matplotlib.pyplot as plt
plt.figure(figsize=(6,4))
plt.plot(results_df['k'], results_df['accuracy'], marker='o')
plt.xlabel("k value")
plt.ylabel("Test Accuracy")
plt.title("KNN Accuracy for Different k (PCA + scaled)")
plt.grid(True)
plt.show()
```



CART Decision Tree

```
In [36]: from sklearn.tree import DecisionTreeClassifier

cart = DecisionTreeClassifier(
    class_weight='balanced',
    random_state=1)

# Train
cart.fit(X_train_tfidf, y_train)

# Predictions
y_pred_cart = cart.predict(X_test_tfidf)

# Metrics
print_classification_metrics(y_test, y_pred_cart,
                             model_name="Decision Tree (CART)")
```

```
===== Decision Tree (CART) =====
Confusion Matrix (rows = true, cols = predicted):
[[1322 555]
 [ 729 6138]]

Accuracy : 0.8532
Precision: 0.9171
Recall   : 0.8938
F1-score : 0.9053
```

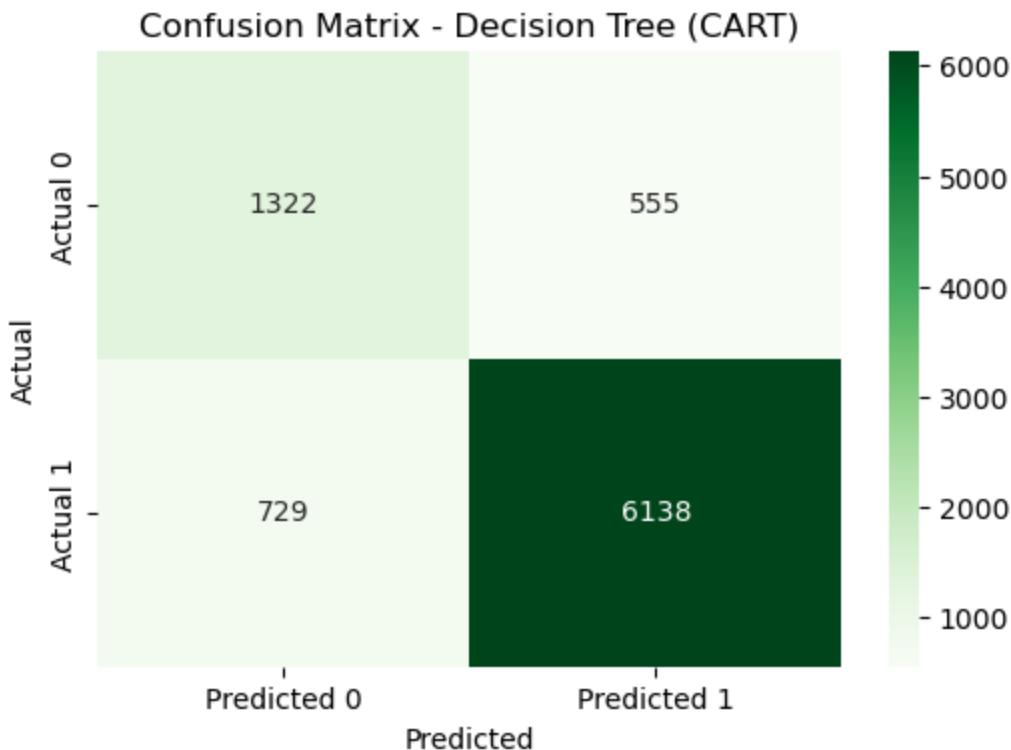
Detailed classification report:

	precision	recall	f1-score	support
0.0	0.6446	0.7043	0.6731	1877
1.0	0.9171	0.8938	0.9053	6867
accuracy			0.8532	8744
macro avg	0.7808	0.7991	0.7892	8744
weighted avg	0.8586	0.8532	0.8555	8744

```
In [37]: y_pred_cart = cart.predict(X_test_tfidf)
acc_cart = accuracy_score(y_test, y_pred_cart)
f1_cart = f1_score(y_test, y_pred_cart)
```

```
In [38]: cm_cart = confusion_matrix(y_test, y_pred_cart)

plt.figure(figsize=(6,4))
sns.heatmap(cm_cart, annot=True, fmt='d', cmap='Greens',
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.title("Confusion Matrix - Decision Tree (CART)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



Naive Bayes

```
In [39]: from sklearn.naive_bayes import MultinomialNB  
  
nb_model = MultinomialNB()  
  
# Train NB on TF-IDF sparse matrix (correct for text)  
nb_model.fit(X_train_tfidf, y_train)  
  
# Predictions  
nb_pred = nb_model.predict(X_test_tfidf)  
  
# Use YOUR metric function here  
print_classification_metrics(y_test, nb_pred,  
                           model_name="Naive Bayes (Multinomial)")
```

```
===== Naive Bayes (Multinomial) =====
Confusion Matrix (rows = true, cols = predicted):
[[1329 548]
 [ 97 6770]]

Accuracy : 0.9262
Precision: 0.9251
Recall   : 0.9859
F1-score : 0.9545

Detailed classification report:
      precision    recall  f1-score   support

          0.0      0.9320    0.7080    0.8047     1877
          1.0      0.9251    0.9859    0.9545     6867

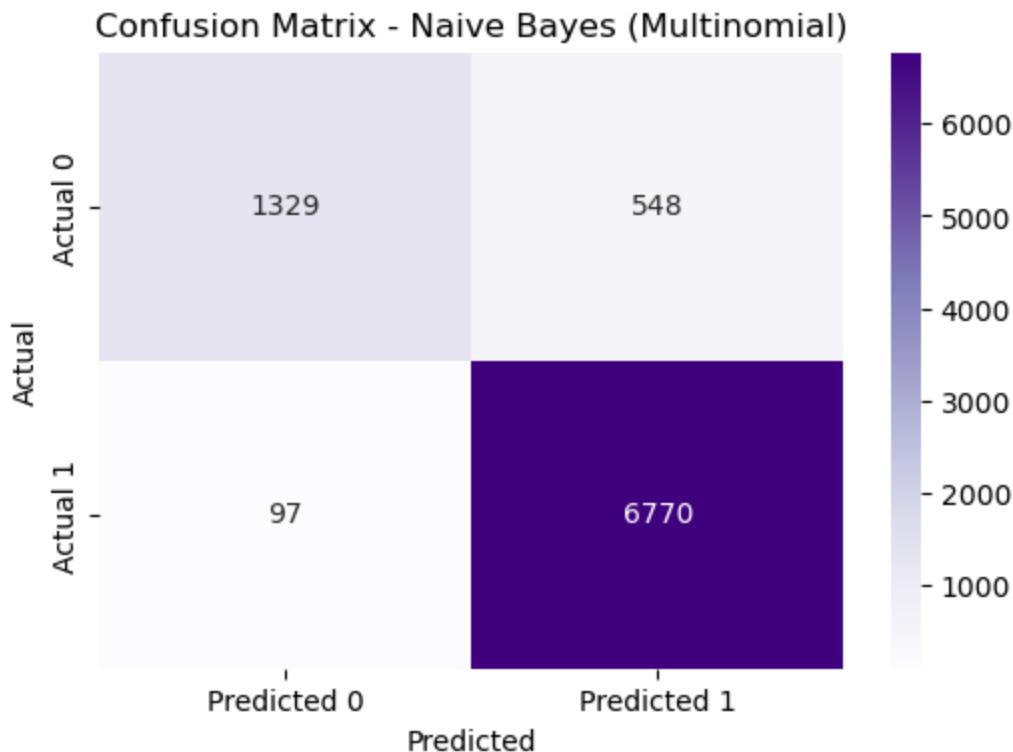
   accuracy                           0.9262     8744
  macro avg       0.9285    0.8470    0.8796     8744
weighted avg     0.9266    0.9262    0.9224     8744
```

```
In [40]: y_pred_nb = nb_model.predict(X_test_tfidf)
acc_nb = accuracy_score(y_test, y_pred_nb)
f1_nb = f1_score(y_test, y_pred_nb)
```

```
In [41]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

cm_nb = confusion_matrix(y_test, nb_pred)

plt.figure(figsize=(6,4))
sns.heatmap(cm_nb, annot=True, fmt='d', cmap='Purples',
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.title("Confusion Matrix - Naive Bayes (Multinomial)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



Multiple Linear Regression

```
In [42]: from sklearn.linear_model import LinearRegression

mlr_model = LinearRegression()
mlr_model.fit(X_train_mlr, y_train)

# Predict continuous values (0-1)
mlr_pred_continuous = mlr_model.predict(X_test_mlr)

# Convert to binary
mlr_pred = (mlr_pred_continuous >= 0.5).astype(int)

# Use your metric function
print_classification_metrics(
    y_test, mlr_pred,
    model_name="Multiple Linear Regression (PCA + scaled)"
)
```

```
===== Multiple Linear Regression (PCA + scaled) =====
Confusion Matrix (rows = true, cols = predicted):
[[1346  531]
 [ 166 6701]]

Accuracy : 0.9203
Precision: 0.9266
Recall   : 0.9758
F1-score : 0.9506

Detailed classification report:
      precision    recall  f1-score   support

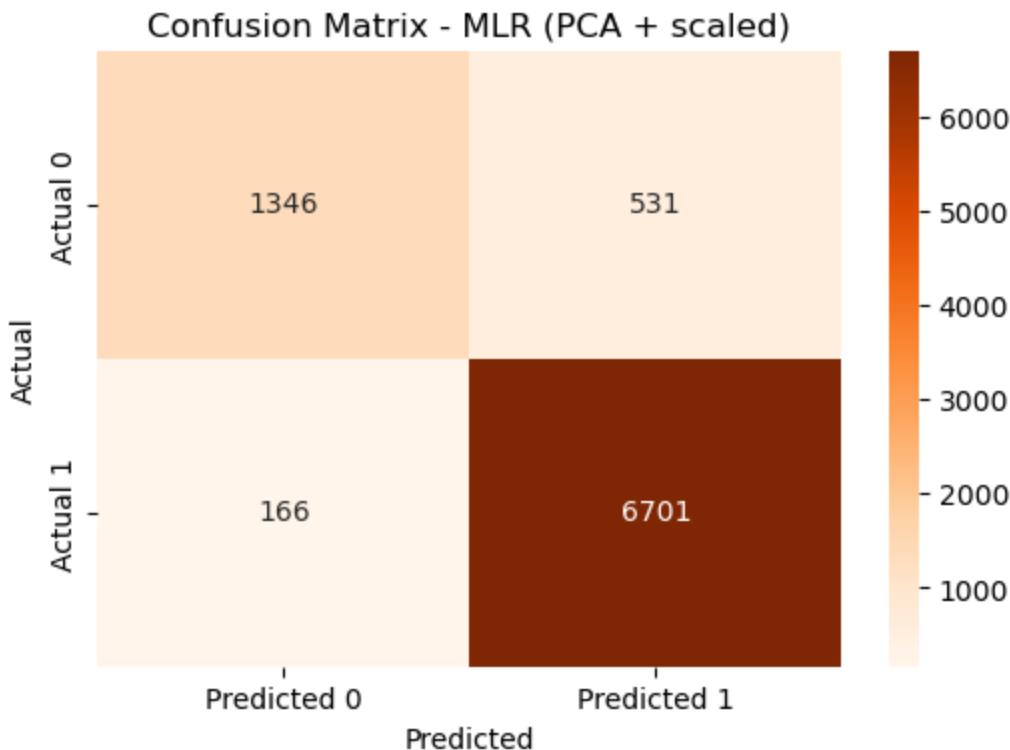
          0.0      0.8902    0.7171    0.7943     1877
          1.0      0.9266    0.9758    0.9506     6867

   accuracy                           0.9203     8744
  macro avg       0.9084    0.8465    0.8724     8744
weighted avg     0.9188    0.9203    0.9170     8744
```

```
In [43]: y_pred_mlr = mlr_pred # already computed
acc_mlr = accuracy_score(y_test, y_pred_mlr)
f1_mlr = f1_score(y_test, y_pred_mlr)
```

```
In [44]: cm_mlr = confusion_matrix(y_test, mlr_pred)

plt.figure(figsize=(6,4))
sns.heatmap(cm_mlr, annot=True, fmt='d', cmap='Oranges',
            xticklabels=['Predicted 0', 'Predicted 1'],
            yticklabels=['Actual 0', 'Actual 1'])
plt.title("Confusion Matrix - MLR (PCA + scaled)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



Results Comparison

```
In [86]: import pandas as pd

comparison_df = pd.DataFrame({
    "Model": ["Logistic Regression", "KNN", "CART", "Naive Bayes", "MLR"],
    "Accuracy": [acc_lr, acc_knn, acc_cart, acc_nb, acc_mlr],
    "F1-score": [f1_lr, f1_knn, f1_cart, f1_nb, f1_mlr]
})

comparison_df
```

```
Out[86]:
```

	Model	Accuracy	F1-score
0	Logistic Regression	0.908394	0.939433
1	KNN	0.846523	0.902044
2	CART	0.853156	0.905310
3	Naive Bayes	0.926235	0.954529
4	MLR	0.920288	0.950564

```
In [46]: plt.figure(figsize=(8,5))
plt.bar(comparison_df["Model"], comparison_df["Accuracy"], color='lightblue')
plt.title("Model Accuracy Comparison")
plt.ylabel("Accuracy")
plt.ylim(0.8, 1.0)
plt.grid(axis='y', linestyle='--', alpha=0.5)
for i, v in enumerate(comparison_df["Accuracy"]):
```

```
plt.text(i, v + 0.002, f'{v:.3f}', ha='center')  
plt.show()
```

